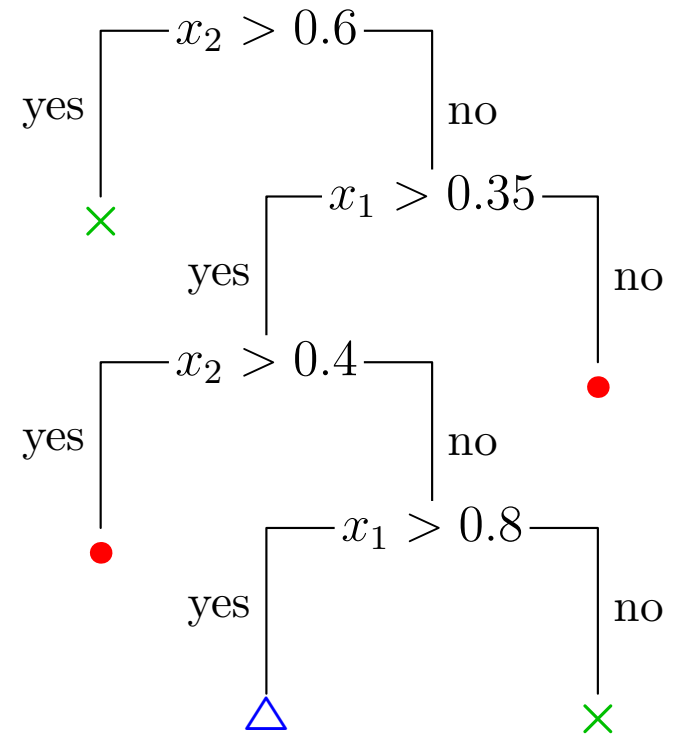
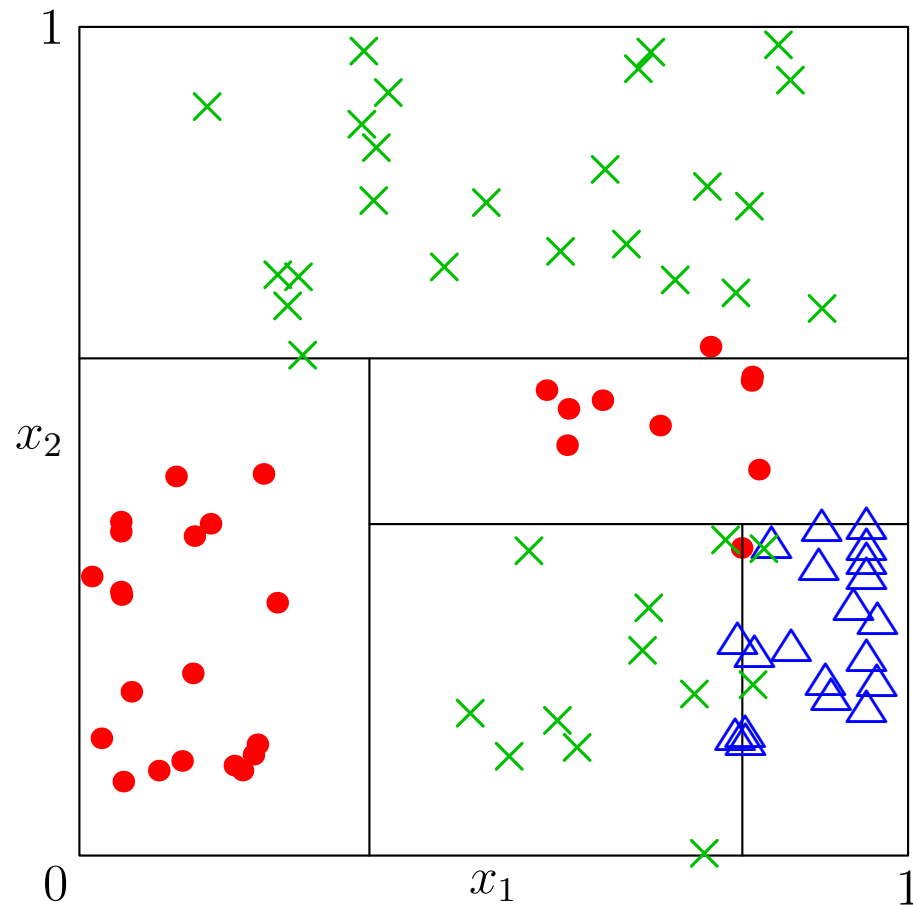


# Advanced Machine Learning

## *Ensemble Methods*

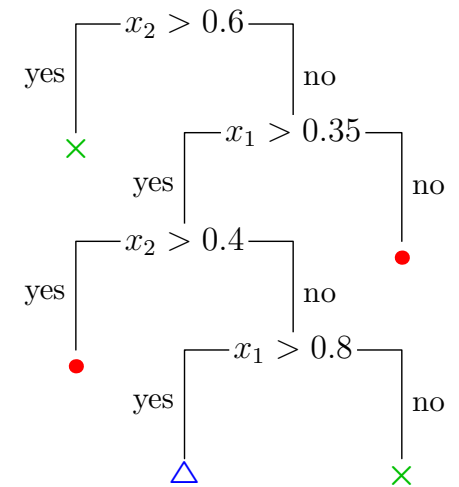
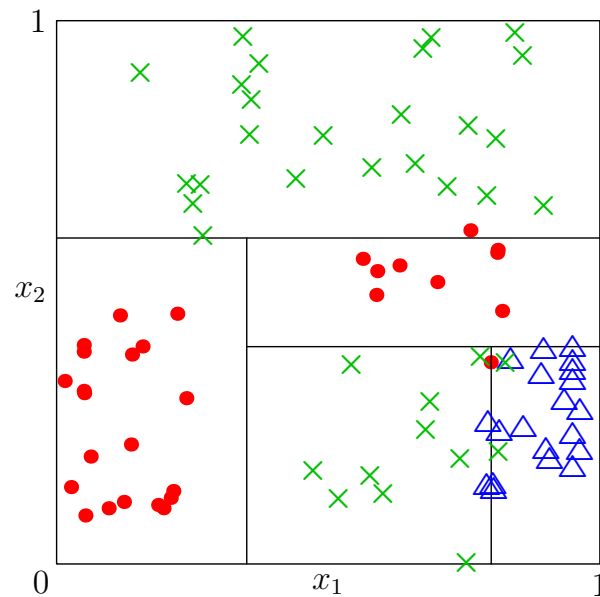


*Decision Trees, Averaging, Bagging*

# Outline

## 1. Decision Trees

## 2. Bagging



# Removing Variance By Averaging

- We can reduce the variance and hence improve our generalisation error by averaging over different learning machines■
- There are a number of different techniques for doing this that go by the name of **ensemble methods** or **ensemble learning**■
- This trick can be used with many different learning machines, but is clearly most practical for machine that can be trained quickly■
- (nevertheless, even for deep learning taking the average response of many machines is usually done to win competitions)■

# Ensembling of Decision Trees

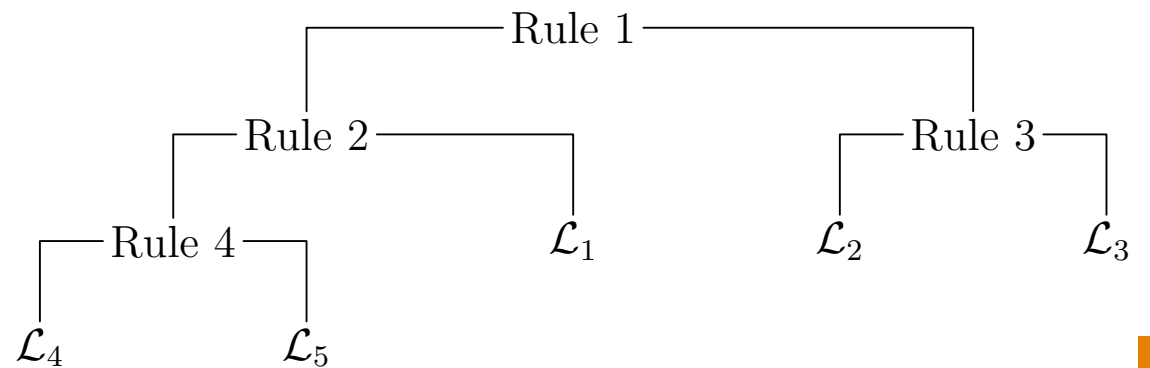
- One set of algorithms where ensembling are common place are decision trees■
- These are particularly good for handling messy data
  - ★ categorical data■
  - ★ mixture of data types■
  - ★ missing data■
  - ★ large data sets■
  - ★ multiclass■
- In many competitions ensembled trees, particularly *random forests* and *gradient boosting* beat all other techniques■

# Decision Trees

- A decision trees builds a binary tree to partition the data,  $\mathcal{D} = \{(\mathbf{x}_i, y_i) | i = 1, \dots, m\}$ , into the leaves of the tree■
- Each decision rule depends on a single feature■
- At each step the rule is chosen that maximise the “*purity*” of the leaf nodes■
- Decisions can be made on numerical values or categories■

# Partitioning

- Consider a classification problems with examples  $(\mathbf{x}, y)$  belonging to some classes  $y \in \mathcal{C}$
- The data is partitioned by the tree into leaves



- The proportion of data points in leaf  $\mathcal{L}$  belonging to class  $c$  is

$$p_c(\mathcal{L}) = \frac{1}{|\mathcal{L}|} \sum_{(\mathbf{x}, y) \in \mathcal{L}} \mathbb{I}[y = c]$$

where  $\mathbb{I}[y = c] = 1$  if  $y = c$  and 0 otherwise

# Leaf Purity



- Two different purity measures,  $Q_m(\mathcal{L})$ , for a leaf node  $\mathcal{L}$  are commonly used

- ★ **Gini index**

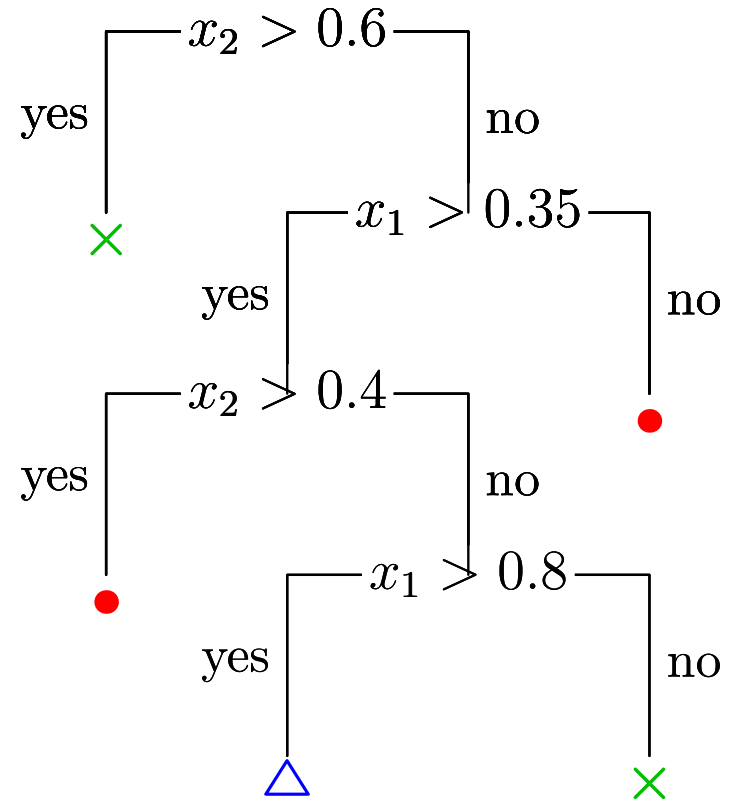
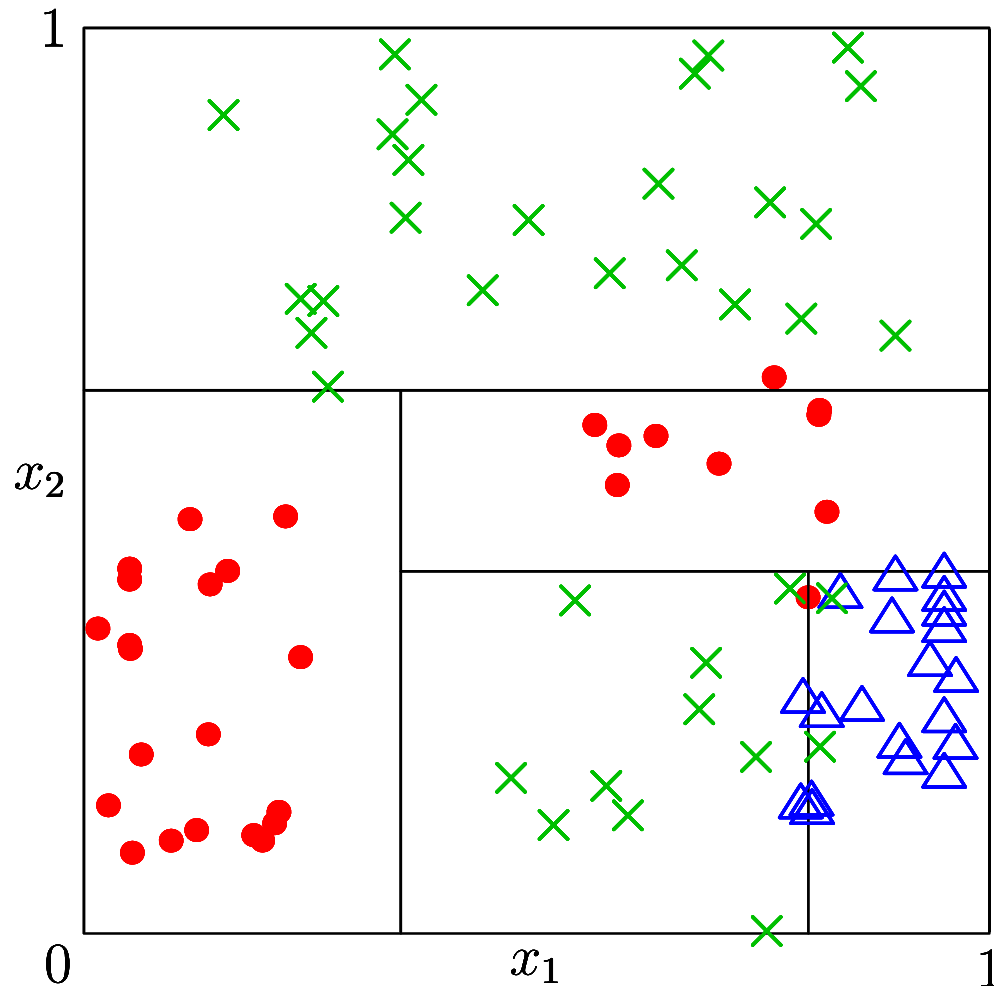
$$Q_m^g(\mathcal{L}) = \sum_{c \in \mathcal{C}} p_c(\mathcal{L}) (1 - p_c(\mathcal{L}))$$

- ★ **Cross-entropy**

$$Q_m^e(\mathcal{L}) = - \sum_{c \in \mathcal{C}} p_c(\mathcal{L}) \log(p_c(\mathcal{L}))$$



# Building Decision Trees





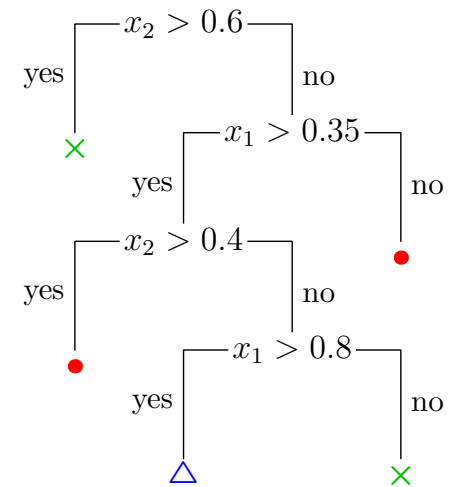
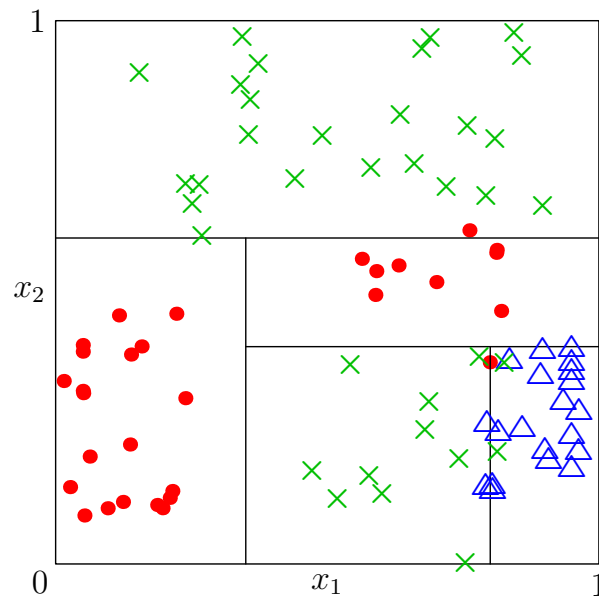
# Observations

- Decision trees are very useful for exploring new data sets—the tree shows what features are most important
- Decision trees can also be used for regression problems
  - ★ Approximate function by a series of rules
  - ★ Reduce variance between data points assigned to leaf nodes
- CART is a classic implementation that builds **C**lassification **A**nd **R**egression **T**rees
- Decision trees depend strongly on the early decisions and so vary a lot for slightly different data sets—high variance

# Outline


## 1. Decision Trees

## 2. Bagging



# Error In The Means

- By taking the mean over many samples we can reduce the variance and thus improve our generalisation performance■
- To get a feel for this consider estimating the mean of a random variable,  $X$ , from a number of samples ( $n = 5$  in the example below)■



# Mean and Variance

- The expected value of the mean,  $\hat{\mu}_n$ , of  $n$  random **independent** variables,  $X_i$ , is the expected value  $\mu = \mathbb{E}[X_i]$


$$\mathbb{E}[\hat{\mu}_n] = \mathbb{E}\left[\frac{1}{n} \sum_{i=1}^n X_i\right] = \frac{1}{n} \sum_{i=1}^n \mathbb{E}[X_i] = \frac{1}{n} \sum_{i=1}^n \mu = \mu$$

- The variance is  $\mathbb{E}[(\hat{\mu}_n - \mu)^2]$  or equivalently

$$\begin{aligned} \frac{1}{n^2} \mathbb{E}\left[\left(\sum_{i=1}^n (X_i - \mu)\right)^2\right] &= \frac{1}{n^2} \mathbb{E}\left[\sum_{i=1}^n (X_i - \mu)^2 + \sum_{i=1}^n \sum_{\substack{j=1 \\ j \neq i}}^n (X_i - \mu)(X_j - \mu)\right] \\ &= \frac{1}{n^2} \sum_{i=1}^n \left( \mathbb{E}[(X_i - \mu)^2] + \sum_{\substack{j=1 \\ j \neq i}}^n \mathbb{E}[X_i - \mu] \mathbb{E}[X_j - \mu] \right) \\ &= \frac{1}{n^2} \sum_{i=1}^n \sigma^2 = \frac{1}{n} \sigma^2 \end{aligned}$$

# Bootstrap Aggregation (Bagging)

- To reduce the variance in a learning machine (such as a decision tree) we can average over many machines■
- To average many machines they must learn something different■
- We only have one data set, but we can resample from the data set to make them look a bit different■—this is known as **bootstrapping**■



$(\mathbf{x}_5, y_5)$   $(\mathbf{x}_4, y_4)$   $(\mathbf{x}_5, y_5)$   $(\mathbf{x}_6, y_6)$   $(\mathbf{x}_2, y_2)$   $(\mathbf{x}_1, y_1)$   
 $(\mathbf{x}_1, y_1)$   $(\mathbf{x}_5, y_5)$   $(\mathbf{x}_3, y_3)$   $(\mathbf{x}_2, y_2)$   $(\mathbf{x}_5, y_5)$   $(\mathbf{x}_2, y_2)$   
 $(\mathbf{x}_6, y_6)$   $(\mathbf{x}_1, y_1)$   $(\mathbf{x}_3, y_3)$   $(\mathbf{x}_6, y_6)$   $(\mathbf{x}_3, y_3)$   $(\mathbf{x}_4, y_4)$   
 $(\mathbf{x}_4, y_4)$   $(\mathbf{x}_3, y_3)$   $(\mathbf{x}_1, y_1)$   $(\mathbf{x}_6, y_6)$   $(\mathbf{x}_4, y_4)$   $(\mathbf{x}_6, y_6)$   
 $(\mathbf{x}_3, y_3)$   $(\mathbf{x}_2, y_2)$   $(\mathbf{x}_3, y_3)$   $(\mathbf{x}_4, y_4)$   $(\mathbf{x}_6, y_6)$   $(\mathbf{x}_2, y_2)$  ■

# Performance of Bagging

- For classification we get our different machines to vote■
- For regression we can average the prediction of different machines■
- Bagging improves the performance of decision trees■
- However, we can usually do better using Boosting■
- This is because our decision trees are correlated■

# Variance of Positive Correlated Variables

- If we calculate the variance of the mean of positively correlated variables with correlation  $\rho$  we find

$$\frac{1}{n^2} \mathbb{E} \left[ \left( \sum_{i=1}^n X_i - n\mu \right)^2 \right] = \rho \sigma^2 + \frac{1-\rho}{n} \sigma^2$$

$$(\rho = \mathbb{E}[(X_i - \mu)(X_j - \mu)] / \sigma^2) \blacksquare$$

- As  $n \rightarrow \infty$  the second term vanishes, but we are left with the first term  $\blacksquare$
- If we want to do well we need our learning machines to be unbiased and decorrelated  $\blacksquare$

# Random Forest

- In random forests we average much less correlated trees■
- To do this for each tree we choose a subset of  $p' \ll p$  of the features on which to split the tree■
- Typically  $p'$  can range from 1 to  $\sqrt{p}$ ■
- The trees aren't that good, but are very decorrelated■
- By averaging over a huge number of trees (order of 1000) we typically get good results■
- Random Forest won (wins?) many competitions■



# Lessons

- Ensemble methods have proved themselves to be very powerful■
- They work by averaging over different machines, trying to reduce their variance■
- Here the variance comes from forcing the machines to learn different functions using Bootstrap Aggregation■
- Tend to work best with very simple models (true of random forest and boosting)■—seems to reduce over-fitting■
- Random forest is very powerful, but gradient boosting is competitive■