


# Advanced Machine Learning


## *Entropy*



## *Entropy, Coding, Maximum Entropy*

# Outline

1. **Measuring Uncertainty**
2. Code Length
3. Maximum Entropy



# Measuring Uncertainty

- What is more uncertain tossing a coin three times or throwing a dice■
- The answer depends on whether you care about the order of the coin tosses■
- But, how do we answer such a question?■
- Let  $X$  be a random variable denoting the possible outcomes■
- Interestingly, Shannon entropy give a precise answer

$$H_X = - \sum_{x \in \mathcal{X}} \mathbb{P}(X = x) \log_2(\mathbb{P}(X = x)) \blacksquare$$

# Let's Calculate

- For an honest dice  $D \in \{1,2,3,4,5,6\}$  and  $\mathbb{P}(D = i) = 1/6$  so

$$H_D = -\sum_{i=1}^6 \frac{1}{6} \log_2 \left( \frac{1}{6} \right) = -\log_2 \left( \frac{1}{6} \right) = \log_2(6) \approx 2.584 \text{bits}$$

- For an honest coin where we care about the order so  $C \in \{000,001,\dots,111\}$  the  $\mathbb{P}(C = i) = \frac{1}{8}$  and

$$H_C = -\sum_{i=0}^7 \frac{1}{8} \log_2 \left( \frac{1}{8} \right) = -\log_2 \left( \frac{1}{8} \right) = \log_2(8) = 3 \text{bits}$$

- This clearly makes sense: there are more possible outcomes; all equally likely

# Unordered Coin Toss

- What if we don't care about the order of the out-come then  $\mathbb{P}(HHH) = \mathbb{P}(TTT) = 1/8$ ,  $\mathbb{P}(HHT) = \mathbb{P}(HTT) = 3/8$  so

$$H_U = -\frac{1}{4}\log_2\left(\frac{1}{8}\right) - \frac{3}{4}\log_2\left(\frac{3}{8}\right) \approx 1.811\text{bits}$$

- This seems reasonable, although it is not obvious how you would determine this without using entropy
- But why Shannon entropy?

# Additive Entropy


- If  $H_X$  and  $H_Y$  is the uncertainty of two independent random variable  $X$  and  $Y$ , what is the uncertainty of the combined event  $(X,Y)$ ?

$$\begin{aligned} H_{(X,Y)} &= - \sum_{X,Y} \mathbb{P}(X,Y) \log_2(\mathbb{P}(X,Y)) \blacksquare \\ &= - \sum_{X,Y} \mathbb{P}(X) \mathbb{P}(Y) \log_2(\mathbb{P}(X) \mathbb{P}(Y)) \blacksquare \\ &= - \sum_{X,Y} \mathbb{P}(X) \mathbb{P}(Y) (\log_2(\mathbb{P}(X)) + \log_2(\mathbb{P}(Y))) \blacksquare \\ &= - \sum_X \mathbb{P}(X) \log_2(\mathbb{P}(X)) - \sum_Y \mathbb{P}(Y) \log_2(Y) \blacksquare = H_X + H_Y \blacksquare \end{aligned}$$

- Shannon's entropy is one of the few functions that satisfy this condition  $\blacksquare$

# Outline

1. Measuring Uncertainty
2. **Code Length**
3. Maximum Entropy



# Why Measure Entropy in Bits

- Suppose we had to communicate a message with  $2^n$  equally likely outputs (e.g. the result of  $n$ -coin tosses)■
- We can do this with a binary string with  $n$  bits (011..0)■
- If there were 5 possible outcomes I could do this with 3 bits, but waste  $3/8$  of the message■
- However if we have a batch of 3 independent messages each with 5 outcomes then there are 125 possible outcomes■ We could communicate this with 8 bits. This would waste  $3/128$  of the message■
- By batching together enough messages with  $N$  outcomes then we asymptotically need just  $\log_2(N)$  bits■



# Different Probabilities

- We “showed” that if we had  $N$  events,  $X_i$ , each with probability,  $\mathbb{P}(X_i) = 1/N$ , we can code the outcomes with a message of length  $-\log_2(\mathbb{P}(X_i)) = \log_2(N)$ ■
- With a shorter message we would not be able to distinguish all possible outcomes from the message■
- What happens if some of outcomes occur with a different probability■

$X_i$ :	1	2	3	4	5	6
$p(X_i)$ :	$\frac{1}{8}$	$\frac{1}{8}$	$\frac{1}{8}$	$\frac{1}{8}$	$\frac{1}{4}$	$\frac{1}{4}$
Code:	000	001	010	011	10	11
$L = -\log_2(p(X_i))$ :	3	3	3	3	2	2

# Shannon's Entropy

- If the probabilities are not equal to  $2^{-n}$  we can still find a code with a length very close to  $-\log_2(\mathbb{P}(X))$  per message by transmitting a large number of messages■
- The length of the message measures the amount of **surprise** on receiving the message■
- Shannon's entropy is the expected length of the message to communicate a random variable  $X$



$$H_X = \mathbb{E}_X[-\log_2(\mathbb{P}(X))] = - \sum_{x \in \mathcal{X}} \mathbb{P}(X = x) \log_2(\mathbb{P}(X = x)) \blacksquare$$

- The expected length is a measure of the uncertainty■(how much information on average we need to convey the outcome)■

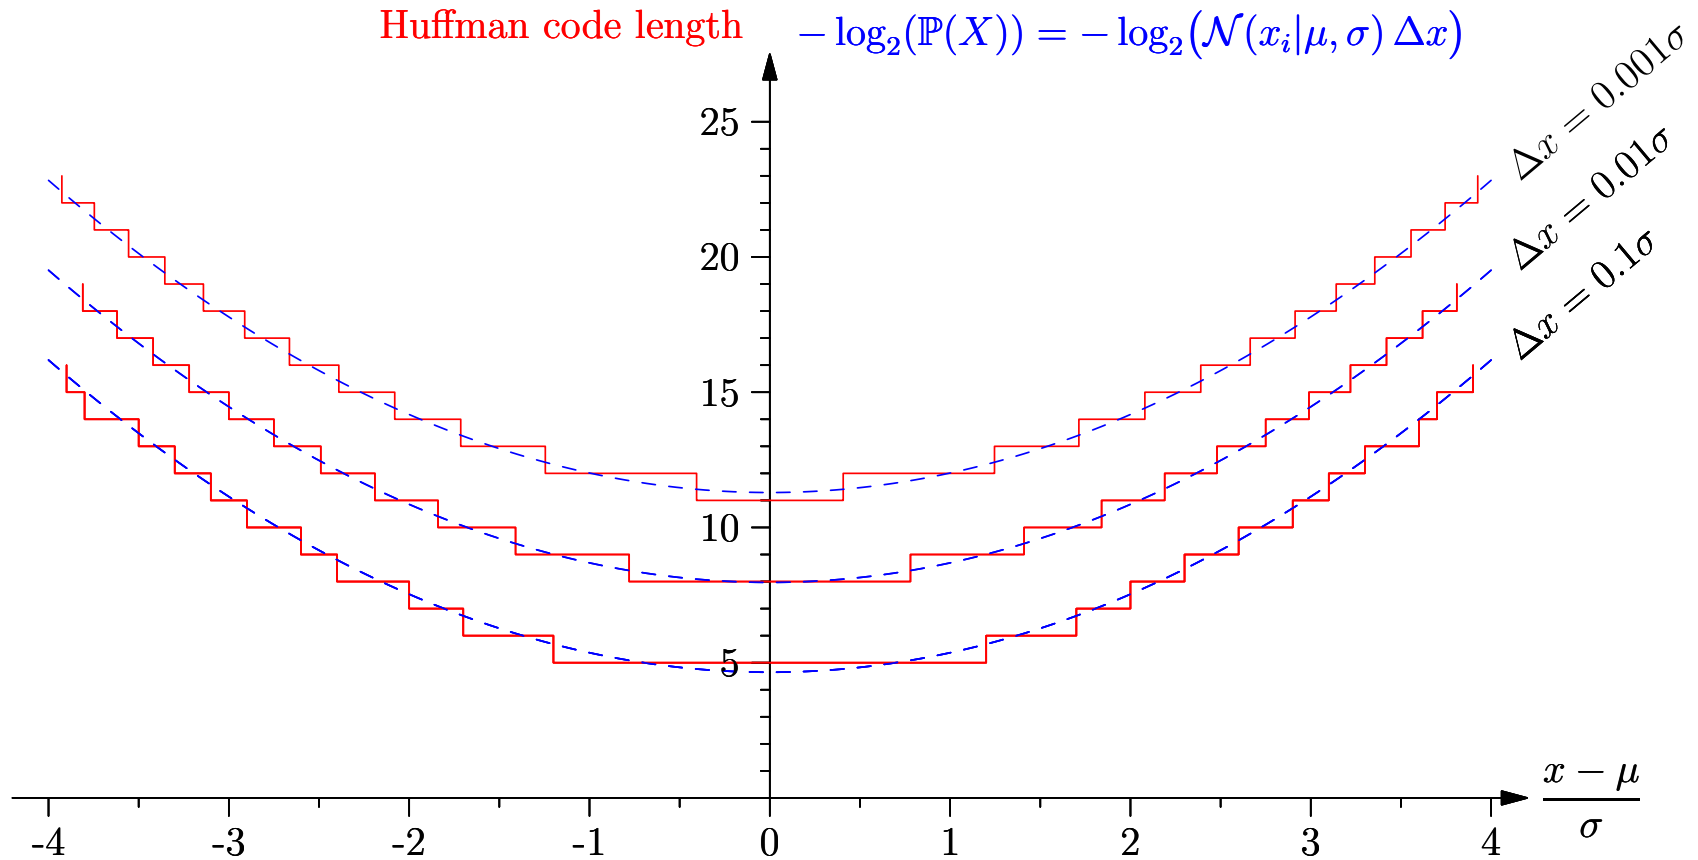
# Real Codes

- Those of you with some computer science background will realise that we can't actually use different length strings in a code without paying some price■
- We won't know where a code word ends so we can't decode the message■
- An optimal solution is to use Huffman encoding where we associate the leaf of a tree with each code word■
- Using the tree we can decode any message constructed using the tree■
- There is a greedy algorithm for constructing the optimal tree■

# Coding Normals



# Coding Normals to Accuracy $\Delta x$



# bits and nats

- We have measured entropy in **bits** using

$$H_X = - \sum_{x \in \mathcal{X}} \mathbb{P}(X = x) \log_2(\mathbb{P}(X = x)) \blacksquare$$


- Sometimes it is easier to use natural logarithms

$$H_X = - \sum_{x \in \mathcal{X}} \mathbb{P}(X = x) \ln(\mathbb{P}(X = x)) \blacksquare$$

- In this case the entropy is measured in **nats** with 1 nat equal to  $\log_2(e)$  bits  $\blacksquare$
- This is often easier when we want to do calculus on entropy  $\blacksquare$


# Outline

1. Measuring Uncertainty
2. Code Length
3. **Maximum Entropy**



# Number of States

- Suppose I have  $N$  balls I then put in  $K$  boxes with coordinates  $x_i$  such that the mean is  $\mu$  and variance is  $\sigma^2$  ■



$$\mathbb{P}(\mathbf{n}) \propto \frac{N!}{n_1!n_2!\cdots n_K!} \left[ \sum_i \frac{n_i}{N} x_i = \mu \right] \left[ \sum_i \frac{n_i}{N} (x_i - \mu)^2 = \sigma^2 \right] \quad \blacksquare$$



# Stirling's Approximation

- We can approximate the factorial  $n!$  using **Stirling's approximation**


$$n! \approx \sqrt{2\pi n} n^n e^{-n}$$

$$\log(n!) = n \log(n) - n + \frac{1}{2} \log(2\pi n) \blacksquare$$

- Using this in our formula for  $\mathbb{P}(\mathbf{n})$  we have

$$\mathbb{P}(\mathbf{n}) \approx C e^{-N \sum_i \frac{n_i}{N} \log\left(\frac{n_i}{N}\right)} \prod_{l=1}^3 \left[ \sum_i \frac{n_i}{N} f_l(x_i) = v_l \right] \blacksquare$$

where  $(f_1(x_i), v_l) = \{(1, 1), (x_i, \mu), ((x_i - \mu)^2, \sigma^2)\} \blacksquare$



# Number of States and Entropy

- Let  $p(x_i) = n_i/N$  be the proportion of balls in bin  $i$  then

$$\mathbb{P}(\mathbf{n}) \approx C e^{N H_X} \prod_{l=1}^3 \left[ \sum_i \frac{n_i}{N} f_l(x_i) = v_l \right]$$

where

$$H_X = - \sum_i p(x_i) \log(p(x_i)) \blacksquare$$

- That is, the “entropy” can be seen as a measure of the logarithm of the number of configurations  $\blacksquare$
- When the number of balls,  $N \rightarrow \infty$  the overwhelmingly likely configurations is the one that maximises the entropy subject to the observed mean and variance  $\blacksquare$

# Maximum Entropy Method

- When we are trying to infer a distribution given some observations then we can maximise the entropy subject to constraints—the entropy acts as a prior
- This is known as the **maximum entropy method**
- We can rationalise this as this is by far the most likely set of configurations consistent with the observations
- Alternatively we can see this as maximising our uncertainty given what we know—being as unbiased as possible
- It only gives a good approximation if all possibilities are equally likely

# Knowing the Mean and Variance

- Consider a continuous random variable,  $X$ , with a known mean and second moment

$$\mathbb{E}[X] = \mu, \quad \mathbb{E}[X^2] = \mu_2 = \mu^2 + \sigma^2$$

- To maximise the entropy subject to constraints consider

$$\begin{aligned} \mathcal{L}(f) = & - \int f_X(x) \log(f_X(x)) dx + \lambda_0 \left( \int f_X(x) dx - 1 \right) \\ & + \lambda_1 \left( \int f_X(x) x dx - \mu \right) + \lambda_2 \left( \int f_X(x) x^2 dx - \mu_2 \right) \end{aligned}$$

- Thus

$$\frac{\delta \mathcal{L}(f)}{\delta f_X(x)} = -\log(f_X(x)) - 1 + \lambda_0 + \lambda_1 x + \lambda_2 x^2 = 0$$

- Or

$$f_X(x) = e^{-1+\lambda_0+\lambda_1 x+\lambda_2 x^2}$$

# Normal Distribution

- We have three constraints

$$\begin{aligned}\int e^{-1+\lambda_0+\lambda_1 x+\lambda_2 x^2} dx &= 1 \\ \int e^{-1+\lambda_0+\lambda_1 x+\lambda_2 x^2} x dx &= \mu \\ \int e^{-1+\lambda_0+\lambda_1 x+\lambda_2 x^2} x^2 dx &= \mu_2 = \mu^2 + \sigma^2\end{aligned}$$

- Solving for  $\lambda_0$ ,  $\lambda_1$  and  $\lambda_2$  then

$$f_X(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-(x-\mu)^2/(2\sigma^2)}$$

- That is, the normal distribution is the maximum entropy distribution given we known the mean and variance

# Using Maximum Entropy

- Maximum entropy is often used to infer distributions■
- It can be very effective, but it might not work well if there are other constraints that we have not included■
- The place that they work superbly well is in statistical physics■
- The whole of statistical physics is about inferring distributions making observations of volume, pressure, etc.■
- Temperature appears rather strangely as a Lagrange multiplier■

# Historic Entropy

- Historically entropy was first introduced in statistical physics by Rudolf Clausius in 1865 (although Macquorn Rankine discussed it in 1850)■
- Its interpretation as the number of states was introduced by Ludwig Boltzmann■
- The person who got it all right was Josiah Willard Gibbs (and James Clerk Maxwell)■
- Claude Shannon invented information theory base on entropy around 1948 (more on that in the next lecture)■
- Ed Jaynes was the first to understand that statistical physics can be seen as an inference problem■

# Conclusion

- Entropy provides a measure of the disorder or uncertainty in a system■
- It forms the basis of information theory which we will look at in the next lecture■
- $-\log(\mathbb{P}(X = x))$  can be seen as the minimum length of a message to communicate  $x$ ■
- This will be used as the basis of the minimum description length formalism also discussed in the next lecture■
- Entropy can be used as a prior, which we often maximise subject to constraints to obtain an unbiased estimate■