


**Entropy***Entropy, Coding, Maximum Entropy***Measuring Uncertainty**

- What is more uncertain tossing a coin three times or throwing a dice?
- The answer depends on whether you care about the order of the coin tosses?
- But, how do we answer such a question?
- Let  $X$  be a random variable denoting the possible outcomes
- Interestingly, Shannon entropy give a precise answer

$$H_X = - \sum_{x \in \mathcal{X}} \mathbb{P}(X = x) \log_2(\mathbb{P}(X = x))$$

**Unordered Coin Toss**


- What if we don't care about the order of the out-come then  $\mathbb{P}(HHH) = \mathbb{P}(TTT) = 1/8$ ,  $\mathbb{P}(HHT) = \mathbb{P}(HTT) = 3/8$  so

$$H_U = -\frac{1}{4} \log_2\left(\frac{1}{8}\right) - \frac{3}{4} \log_2\left(\frac{3}{8}\right) \approx 1.811 \text{ bits}$$


- This seems reasonable, although it is not obvious how you would determine this without using entropy
- But why Shannon entropy?

**Outline**

1. Measuring Uncertainty
2. Code Length
3. Maximum Entropy



1. Measuring Uncertainty
2. Code Length
3. Maximum Entropy

**Let's Calculate**

- For an honest dice  $D \in \{1, 2, 3, 4, 5, 6\}$  and  $\mathbb{P}(D = i) = 1/6$  so

$$H_D = - \sum_{i=1}^6 \frac{1}{6} \log_2\left(\frac{1}{6}\right) = -\log_2\left(\frac{1}{6}\right) = \log_2(6) \approx 2.584 \text{ bits}$$

- For an honest coin where we care about the order so  $C \in \{000, 001, \dots, 111\}$  the  $\mathbb{P}(C = i) = \frac{1}{8}$  and

$$H_C = - \sum_{i=0}^7 \frac{1}{8} \log_2\left(\frac{1}{8}\right) = -\log_2\left(\frac{1}{8}\right) = \log_2(8) = 3 \text{ bits}$$

- This clearly makes sense: there are more possible outcomes; all equally likely

**Additive Entropy**

- If  $H_X$  and  $H_Y$  is the uncertainty of two independent random variable  $X$  and  $Y$ , what is the uncertainty of the combined event  $(X, Y)$ ?

$$\begin{aligned} H_{(X,Y)} &= - \sum_{X,Y} \mathbb{P}(X,Y) \log_2(\mathbb{P}(X,Y)) \\ &= - \sum_{X,Y} \mathbb{P}(X) \mathbb{P}(Y) \log_2(\mathbb{P}(X) \mathbb{P}(Y)) \\ &= - \sum_{X,Y} \mathbb{P}(X) \mathbb{P}(Y) (\log_2(\mathbb{P}(X)) + \log_2(\mathbb{P}(Y))) \\ &= - \sum_X \mathbb{P}(X) \log_2(\mathbb{P}(X)) - \sum_Y \mathbb{P}(Y) \log_2(\mathbb{P}(Y)) = H_X + H_Y \end{aligned}$$

- Shannon's entropy is one of the few functions that satisfy this condition

**Why Measure Entropy in Bits**

- Suppose we had to communicate a message with  $2^n$  equally likely outputs (e.g. the result of  $n$ -coin tosses)
- We can do this with a binary string with  $n$  bits (011..0)
- If there were 5 possible outcomes I could do this with 3 bits, but waste  $3/8$  of the message
- However if we have a batch of 3 independent messages each with 5 outcomes then there are 125 possible outcomes. We could communicate this with 8 bits. This would waste  $3/128$  of the message
- By batching together enough messages with  $N$  outcomes then we asymptotically need just  $\log_2(N)$  bits

## Different Probabilities


- We “showed” that if we had  $N$  events,  $X_i$ , each with probability,  $\mathbb{P}(X_i) = 1/N$ , we can code the outcomes with a message of length  $-\log_2(\mathbb{P}(X_i)) = \log_2(N)$
- With a shorter message we would not be able to distinguish all possible outcomes from the message
- What happens if some of outcomes occur with a different probability

|                         |               |               |               |               |               |               |
|-------------------------|---------------|---------------|---------------|---------------|---------------|---------------|
| $X_i$ :                 | 1             | 2             | 3             | 4             | 5             | 6             |
| $p(X_i)$ :              | $\frac{1}{8}$ | $\frac{1}{8}$ | $\frac{1}{8}$ | $\frac{1}{8}$ | $\frac{1}{4}$ | $\frac{1}{4}$ |
| Code:                   | 000           | 001           | 010           | 011           | 10            | 11            |
| $L = -\log_2(p(X_i))$ : | 3             | 3             | 3             | 3             | 2             | 2             |

## Real Codes


- Those of you with some computer science background will realise that we can't actually use different length strings in a code without paying some price
- We won't know where a code word ends so we can't decode the message
- An optimal solution is to use Huffman encoding where we associate the leaf of a tree with each code word
- Using the tree we can decode any message constructed using the tree
- There is a greedy algorithm for constructing the optimal tree

## Coding Normals to Accuracy $\Delta x$



## Outline

- Measuring Uncertainty
- Code Length
- Maximum Entropy




## Shannon's Entropy

- If the probabilities are not equal to  $2^{-n}$  we can still find a code with a length very close to  $-\log_2(\mathbb{P}(X))$  per message by transmitting a large number of messages
- The length of the message measures the amount of **surprise** on receiving the message
- Shannon's entropy is the expected length of the message to communicate a random variable  $X$

$$H_X = \mathbb{E}_X[-\log_2(\mathbb{P}(X))] = -\sum_{x \in \mathcal{X}} \mathbb{P}(X = x) \log_2(\mathbb{P}(X = x))$$

- The expected length is a measure of the uncertainty (how much information on average we need to convey the outcome)

## Coding Normals



## bits and nats

- We have measured entropy in **bits** using

$$H_X = -\sum_{x \in \mathcal{X}} \mathbb{P}(X = x) \log_2(\mathbb{P}(X = x))$$


- Sometimes it is easier to use natural logarithms

$$H_X = -\sum_{x \in \mathcal{X}} \mathbb{P}(X = x) \ln(\mathbb{P}(X = x))$$

- In this case the entropy is measured in **nats** with 1 nat equal to  $\log_2(e)$  bits
- This is often easier when we want to do calculus on entropy

## Number of States

- Suppose I have  $N$  balls I then put in  $K$  boxes so that  $n_i$  balls finish up in box  $i$



- Counting the number of ways to do this

$$\mathbb{P}(\mathbf{n}) = \frac{N!}{n_1! n_2! \dots n_K!}$$

## Stirling's Approximation

- We can approximate the factorial  $n!$  using **Stirling's approximation**


$$n! \approx \sqrt{2\pi n} n^n e^{-n}$$

$$\log(n!) = n \log(n) - n + \frac{1}{2} \log(2\pi n)$$

- Using this in our formula for  $\mathbb{P}(\mathbf{n})$  we have

$$\mathbb{P}(\mathbf{n}) = \frac{N!}{n_1! n_2! \dots n_K!} = e^{\log(N!) - \sum_i \log(n_i!)} \\ \approx C e^{N \log(N) - \sum_i n_i \log(n_i)} \approx C e^{-N \sum_i \frac{n_i}{N} \log\left(\frac{n_i}{N}\right)}$$

Order  $n$  terms cancel as  $N = \sum_i n_i$



## Maximum Entropy Method

- When we are trying to infer a distribution given some observations then we can maximise the entropy subject to constraints—the entropy acts as a prior
- This is known as the **maximum entropy method**
- We can rationalise this as this is by far the most likely set of configurations consistent with the observations
- Alternatively we can see this as maximising our uncertainty given what we know—being as unbiased as possible
- It only gives a good approximation if all possibilities are equally likely

## Normal Distribution

- We have three constraints

$$\int e^{-1+\lambda_0+\lambda_1 x+\lambda_2 x^2} dx = 1$$

$$\int e^{-1+\lambda_0+\lambda_1 x+\lambda_2 x^2} x dx = \mu$$

$$\int e^{-1+\lambda_0+\lambda_1 x+\lambda_2 x^2} x^2 dx = \mu^2 + \sigma^2$$

- Solving for  $\lambda_0$ ,  $\lambda_1$  and  $\lambda_2$  then

$$f_X(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-(x-\mu)^2/(2\sigma^2)}$$

- That is, the normal distribution is the maximum entropy distribution given we know the mean and variance

## Historic Entropy

- Historically entropy was first introduced in statistical physics by Rudolf Clausius in 1865 (although Macquorn Rankine discussed it in 1850)
- Its interpretation as the number of states was introduced by Ludwig Boltzmann
- The person who got it all right was Josiah Willard Gibbs (and James Clerk Maxwell)
- Claude Shannon invented information theory based on entropy around 1948 (more on that in the next lecture)
- Ed Jaynes was the first to understand that statistical physics can be seen as an inference problem

## Number of States and Entropy

- Let  $p(x_i) = n_i/N$  be the proportion of balls in bin  $i$  then

$$\mathbb{P}(\mathbf{n}) \approx C e^{N H_X}$$

where

$$H_X = -\sum_i p(x_i) \log(p(x_i))$$

- That is, the “entropy” can be seen as a measure of the logarithm of the number of configurations
- When the number of balls,  $N \rightarrow \infty$  the overwhelmingly likely configurations is the one that maximises the entropy subject to the observed mean and variance

## Knowing the Mean and Variance

- Consider a continuous random variable,  $X$ , with a known mean and second moment

$$\mathbb{E}[X] = \mu, \quad \mathbb{E}[X^2] = \mu^2 + \sigma^2$$

- To maximise the entropy subject to constraints consider

$$\mathcal{L}(f) = -\int f_X(x) \log(f_X(x)) dx + \lambda_0 \left( \int f_X(x) dx - 1 \right) \\ + \lambda_1 \left( \int f_X(x) x dx - \mu \right) + \lambda_2 \left( \int f_X(x) x^2 dx - \mu^2 - \sigma^2 \right)$$

- Thus

$$\frac{\delta \mathcal{L}(f)}{\delta f_X(x)} = -\log(f_X(x)) - 1 + \lambda_0 + \lambda_1 x + \lambda_2 x^2 = 0$$

- Or

$$f_X(x) = e^{-1+\lambda_0+\lambda_1 x+\lambda_2 x^2}$$

## Using Maximum Entropy

- Maximum entropy is often used to infer distributions
- It can be very effective, but it might not work well if there are other constraints that we have not included
- The place that they work superbly well is in statistical physics
- The whole of statistical physics is about inferring distributions making observations of volume, pressure, etc.
- Temperature appears rather strangely as a Lagrange multiplier

## Conclusion

- Entropy provides a measure of the disorder or uncertainty in a system
- It forms the basis of information theory which we will look at in the next lecture
- $-\log(\mathbb{P}(X = x))$  can be seen as the minimum length of a message to communication  $x$
- This will be used as the basis of the minimum description length formalism also discussed in the next lecture
- Entropy can be used as a prior, which we often maximise subject to constraints to obtain an unbiased estimate