


Advanced Machine Learning


Understand Mappings



Mappings, Linear Maps, Solving Linear Systems


Outline

1. **Mappings**
2. Linear Maps



Transforming Data

- In the last lecture we spent time developing a sophisticated view of vector spaces and operators
- At a mathematical level machine learning can be viewed as performing an inverse mapping



- Although our mappings are not necessarily linear in either direction we learn a lot by understanding linear operators

Inverse Problems


- Given m observations $\{(\mathbf{x}_k, y_k) | k = 1, \dots, m\}$ and p unknown $\mathbf{w} = (w_1, w_2, \dots, w_p)^\top$ such that $\mathbf{x}_k^\top \mathbf{w} = y_k$ then to find $\mathbf{w}$ ■
- Define the *design matrix* as the matrix of feature vectors

$$\mathbf{X} = \begin{pmatrix} \mathbf{x}_1^\top \\ \mathbf{x}_2^\top \\ \vdots \\ \mathbf{x}_m^\top \end{pmatrix} = \begin{pmatrix} x_{11} & x_{12} & \cdots & x_{1p} \\ x_{21} & x_{22} & \cdots & x_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ x_{m1} & x_{m2} & \cdots & x_{mp} \end{pmatrix} \quad \blacksquare$$

- and the target vector $\mathbf{y} = (y_1, y_2, \dots, y_m)^\top$ ■
- Then if $m = p$ we have $\mathbf{y} = \mathbf{X}\mathbf{w}$ or $\mathbf{w} = \mathbf{X}^{-1}\mathbf{y}$ ■

Linear Regression

- $x_k^T w$ depends on distance from separating plane



- If $m > p$ then X isn't square so doesn't have an inverse
- Worse unless the data is accurate $y \approx Xw \Rightarrow$ no "solution"
- Problem solved by Gauss to predict the orbit of the asteroid Ceres

Linear Least Squares

- The error of input pattern $\mathbf{x}_k$ is

$$\epsilon_k = \mathbf{x}_k^\top \mathbf{w} - y_k$$

- The squared error

$$E(\mathbf{w}|\mathcal{D}) = \sum_{k=1}^m (\mathbf{x}_k^\top \mathbf{w} - y_k)^2 = \sum_{k=1}^m \epsilon_k^2 = \|\boldsymbol{\epsilon}\|^2$$

- We can define the error vector


$$\boldsymbol{\epsilon} = \mathbf{X}\mathbf{w} - \mathbf{y}$$

(note that $\epsilon_k = \mathbf{x}_k^\top \mathbf{w} - y_k$)

- Minimising this error is known as the least squares problem

Finding a Minimum

- The minima of a one dimensional function, $f(x)$, are given by $f'(x) = 0$



- The minima of an n -dimensions function $f(\mathbf{x})$ are given by the set of equations

$$\frac{\partial f(\mathbf{x})}{\partial x_i} = 0 \quad \forall i = 1, \dots, n$$

Gradients

- The **grad** operator ∇ is the gradient operator in high dimensions

$$\nabla f(\mathbf{x}) = \begin{pmatrix} \frac{\partial f(\mathbf{x})}{\partial x_1} \\ \frac{\partial f(\mathbf{x})}{\partial x_2} \\ \vdots \\ \frac{\partial f(\mathbf{x})}{\partial x_n} \end{pmatrix} \quad \blacksquare$$

- The partial derivatives (curly d's)

$$\frac{\partial f(\mathbf{x})}{\partial x_i}$$

means differentiate with respect to x_i treating all other components x_j as constants $\blacksquare$

Least Squares Solution

- The least squared solution is give by

$$\begin{aligned}\nabla E(\boldsymbol{w}|\mathcal{D}) &= \nabla \|\boldsymbol{\epsilon}\|^2 = \nabla \|\boldsymbol{X}\boldsymbol{w} - \boldsymbol{y}\|^2 \\ &= \nabla (\boldsymbol{w}^\top \boldsymbol{X}^\top \boldsymbol{X} \boldsymbol{w} - 2\boldsymbol{w}^\top \boldsymbol{X}^\top \boldsymbol{y} + \boldsymbol{y}^\top \boldsymbol{y}) \\ &= 2(\boldsymbol{X}^\top \boldsymbol{X} \boldsymbol{w} - \boldsymbol{X}^\top \boldsymbol{y}) = 0\end{aligned}$$

- Or

$$\boldsymbol{w} = (\boldsymbol{X}^\top \boldsymbol{X})^{-1} \boldsymbol{X}^\top \boldsymbol{y} = \boldsymbol{X}^+ \boldsymbol{y}$$

- $\boldsymbol{X}^+ = (\boldsymbol{X}^\top \boldsymbol{X})^{-1} \boldsymbol{X}^\top$ is known as the pseudo inverse

Missing Bits of the Mathematics

- Note that $\|a\|^2 = a^\top a = \sum_i a_i^2$

$$\begin{aligned}\|Xw - y\|^2 &= (Xw - y)^\top (Xw - y) = (w^\top X^\top - y^\top)(Xw - y) \\ &= w^\top X^\top Xw - 2w^\top X^\top y + y^\top y\end{aligned}$$

- Where we have used $w^\top X^\top y = y^\top Xw$, $\sum_{i,j} w_i X_{ji} y_j = \sum_{i,j} y_i X_{ij} w_j$
- Also $\nabla w^\top M w = M w + M^\top w$
- If $M = M^\top$ (i.e. M is symmetric) then $\nabla w^\top M w = 2M w$
- As $(AB)^\top = B^\top A^\top$ then $(X^\top X)^\top = X^\top X^{\top\top} = X^\top X$ and $X^\top X$ is symmetric

Computing Gradients


- To understand gradients we sometimes need to go back to components

$$\begin{aligned}\nabla \mathbf{w}^\top \mathbf{M} \mathbf{w} &= \begin{pmatrix} \frac{\partial}{\partial w_1} \\ \frac{\partial}{\partial w_2} \\ \frac{\partial}{\partial w_3} \\ \vdots \end{pmatrix} \sum_{i,j} w_i M_{ij} w_j = \begin{pmatrix} \sum_j M_{1j} w_j + \sum_i w_i M_{i1} \\ \sum_j M_{2j} w_j + \sum_i w_i M_{i2} \\ \sum_j M_{3j} w_j + \sum_i w_i M_{i3} \\ \vdots \end{pmatrix} \\ &= \mathbf{M} \mathbf{w} + \mathbf{M}^\top \mathbf{w}\end{aligned}$$

- It is tedious to compute these things component-wise, but when you need to understand what is going on then go back to the basics

Outline

1. Mappings
2. **Linear Maps**




Solving Inverse Problems

- Gauss showed us how to solve **over-constrained** problems (we have more observations than parameters)■
- We seek a solution which isn't necessarily exact but minimises an error■
- But, what if we have more parameters than observations■
- That is, we are **under-constrained**■
- Note that in some directions you might be over-constrained and in other directions under-constrained■
- This is very typical of most machine learning problems■


Under Constrained Systems

- If we have less data-points than parameters then there will be multiple solutions



What is the Inverse?

- Many points can map to the same points




Under-constrained Systems

- The system is **under-constrained**
- We have more unknowns than equations
- The inverse is not unique
- Solving the inverse problem ($w = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$) is said to be **ill-posed**
- The inverse $(\mathbf{X}^T \mathbf{X})^{-1}$ doesn't exist
- If we have a complicated learning machine and not sufficient data we often end with an ill-posed inverse problem (there are lots of sets of parameters that explain the data)

Ill-Conditions

- Singular matrices are rare (although they occur when we don't have enough data), but matrices that are close to being singular are common■
- If a matrix is close to singular it is ill-conditioned■
- Ill-conditioned matrices have some small eigenvalues■
- All points get contracted towards a plane■
- Large matrices are very often ill conditioned■

III-Conditioned Matrices



III-Conditioning in ML

- Ill-conditioning in machine learning occurs when a very small change in the learning data causes a large change in the predictions of the learning machine■
- In linear regression the matrix $\mathbf{X}^T\mathbf{X}$ is ill-conditioned when we have as many data points as parameters■
- Much of machine learning is concerned with making learning machines better conditioned■
- Adding regularisers is one approach to achieve this■

Summary

- Linear mappings are commonly used in machine learning algorithms such as regression■
- We will often meet the pseudo-inverse involving inverting $\mathbf{X}^T \mathbf{X}$ ■
- They can be inherently unstable to noise in the inputs■