

Support Vector Machines, maximum margins

Support Vector Machines

- Support vector machines, when used right, often have the best generalisation results
- They are typically used on numerical data, but can and have been adapted to text, sequences, etc.
- Although not as trendy as deep learning, they will often be the method of choice on small data sets
- They subtly regularise themselves, choosing a solution that generalises well from a host of different solutions

Extended Feature Space

- To increase the likelihood of linear-separability we often use a high-dimensional mapping

$$\mathbf{x} = (x_1, x_2, \dots, x_p)^T \rightarrow \phi(\mathbf{x}) = (\phi_1(\mathbf{x}), \phi_2(\mathbf{x}), \dots, \phi_r(\mathbf{x}))^T$$

$$r \gg p$$

- Finding the maximum margin hyper-plane is time consuming in “primal” form if r is large
- We can work in the “dual” space of patterns, then we only need to compute inner-products

$$\langle \phi(\mathbf{x}_i), \phi(\mathbf{x}_j) \rangle = \phi(\mathbf{x}_i)^T \phi(\mathbf{x}_j) = \sum_{k=1}^r \phi_k(\mathbf{x}_i) \phi_k(\mathbf{x}_j)$$

Kernel Functions

- Kernel functions are symmetric functions of two variable
- Strong restriction: *positive semi-definite*
- Examples


Quadratic kernel: $K(\mathbf{x}_1, \mathbf{x}_2) = (\mathbf{x}_1^T \mathbf{x}_2)^2$

Gaussian (RBF) kernel: $K(\mathbf{x}_1, \mathbf{x}_2) = e^{-\gamma \|\mathbf{x}_1 - \mathbf{x}_2\|^2}$

- Consider the mapping


$$\mathbf{x}_i = \begin{pmatrix} x_i \\ y_i \end{pmatrix} \rightarrow \phi(\mathbf{x}_i) = \begin{pmatrix} x_i^2 \\ y_i^2 \\ \sqrt{2}x_i y_i \end{pmatrix}$$

- The Big Picture
- Maximum Margins
- Duality
- Practice



Linear Separation of Data

- SVMs classify linearly separable data



- Finds maximum-margin separating plane

Kernel Trick

- If we choose a **positive semi-definite** kernel function $K(\mathbf{x}, \mathbf{y})$ then there exists functions $\phi(\mathbf{x}) = (\phi_k(\mathbf{x}) | k = 1, 2, \dots, r)$, such that

$$K(\mathbf{x}_i, \mathbf{x}_j) = \langle \phi(\mathbf{x}_i), \phi(\mathbf{x}_j) \rangle$$


(like an eigenvector decomposition of a matrix)

- Never need to compute $\phi_k(\mathbf{x}_i)$ explicitly as we only need the inner-product $\langle \phi(\mathbf{x}_i), \phi(\mathbf{x}_j) \rangle = K(\mathbf{x}_i, \mathbf{x}_j)$ to compute maximum margin separating hyper-plane
- Sometimes $\phi(\mathbf{x}_i)$ is an infinite dimensional vector so it is good we don't have to compute all the elements!

Non-linear Separation of Data


$$K(\mathbf{x}_1, \mathbf{x}_2) = \begin{pmatrix} x_1^2 & y_1^2 & \sqrt{2}x_1y_1 \end{pmatrix} \begin{pmatrix} x_2^2 \\ y_2^2 \\ \sqrt{2}x_2y_2 \end{pmatrix} = x_1^2x_2^2 + y_1^2y_2^2 + 2x_1y_1x_2y_2$$

$$= (x_1x_2 + y_1y_2)^2 = (\mathbf{x}_1^T \mathbf{x}_2)^2$$



Outline

1. The Big Picture
2. Maximum Margins
3. Duality
4. Practice




Inner Product

- Recall the inner or dot product in \mathbb{R}^n

$$\langle x, y \rangle = x \cdot y = x^T y = \sum_{i=1}^n x_i y_i = \|x\| \|y\| \cos(\theta)$$

- If $\|w\| = 1$ then $\langle x, w \rangle = \|x\| \cos(\theta)$




Maximise Margin

- Consider a linearly separable set of data
 - $\mathcal{D} = \{(x_k, y_k)\}_{k=1}^m$
 - $y_k \in \{-1, 1\}$
- Our task is to find a separating plane defined by the orthogonal vector w and a threshold b such that

$$y_k \left(\frac{\langle w, x_k \rangle}{\|w\|} - b \right) \geq \Delta$$

where Δ is the margin

Distance to hyperplanes



Constrained Optimisation

- Wish to find w and b to maximise Δ subject to constraints

$$y_k \left(\frac{\langle w, x_k \rangle}{\|w\|} - b \right) \geq \Delta \quad \text{for all } k = 1, 2, \dots, m$$

- If we divide through by Δ

$$y_k \left(\frac{\langle w, x_k \rangle}{\Delta \|w\|} - \frac{b}{\Delta} \right) \geq 1 \quad \text{for all } k = 1, 2, \dots, m$$

- Define $\hat{w} = w / (\Delta \|w\|)$ and $\hat{b} = b / \Delta$

$$y_k \left(\langle \hat{w}, x_k \rangle - \hat{b} \right) \geq 1$$

Quadratic Programming Problem


- Note that as $\hat{w} = w / (\Delta \|w\|)$

$$\|\hat{w}\| = \left\| \frac{w}{\Delta \|w\|} \right\| = \frac{1}{\Delta \|w\|} \|w\| = \frac{1}{\Delta}$$

- Minimising $\|\hat{w}\|^2$ is equivalent to maximising the margin Δ
- Can write the optimisation problem as a *quadratic programming problem*

$$\min_{\hat{w}, \hat{b}} \frac{\|\hat{w}\|^2}{2} \quad \text{subject to } y_k \left(\langle \hat{w}, x_k \rangle - \hat{b} \right) \geq 1 \quad \text{for all } k = 1, 2, \dots, m$$

Quadratic Programming in SVMs




Quadratic Programming

- We have a quadratic programming problem for the weights $\hat{w} = (\hat{w}_1, \hat{w}_2, \dots, \hat{w}_p)$ and bias \hat{b} and m constraints
- This is a classic but fiddly optimisation problems
- It can be solved in $O(p^3)$ time (it involves inverting matrices) (pew it is not NP-complete!)
- We will see that there is an equivalent dual problem which allows us to use the kernel trick with time complexity $O(m^3)$

Outline

1. The Big Picture
2. Maximum Margins
3. Duality
4. Practice



Lagrangian

- In the extended feature space we can find a separating plane (given by w and b) with maximum margin by solving the problem

$$\min_{w,b} \frac{\|w\|^2}{2} \quad \text{subject to } y_k(\langle w, \phi(x_k) \rangle - b) \geq 1 \text{ for all } k = 1, 2, \dots, m$$

- We can write this as a Lagrange problem

$$\min_{w,b} \max_{\alpha \geq 0} \mathcal{L}(w, b, \alpha)$$

where

$$\mathcal{L}(w, b, \alpha) = \frac{\|w\|^2}{2} - \sum_{k=1}^m \alpha_k (y_k (\langle w, \phi(x_k) \rangle - b) - 1)$$

subject to $\alpha_k \geq 0$

The Dual Problem

- The dual problem is now to find α_k 's that maximise

$$\mathcal{L}(\alpha) = \sum_{k=1}^m \alpha_k - \frac{1}{2} \sum_{k,l=1}^m \alpha_k \alpha_l y_k y_l \langle \phi(x_k), \phi(x_l) \rangle$$

- subject to constraints

$$\sum_{k=1}^m \alpha_k y_k = 0 \quad \forall k = 1, 2, \dots, m \quad \alpha_k \geq 0$$

- The Hessian of $\mathcal{L}(\alpha)$ has elements $H_{kl} = -y_k y_l \langle \phi(x_k), \phi(x_l) \rangle$ so $v^T H v = -\|\sum_k v_k y_k \phi(x_k)\|^2 \leq 0$ (note this is negative semi-definite so there is a unique maximum)

Sequential Minimal Optimisation

- One of the most efficient techniques for training SVMs is *Sequential Minimal Optimisation* or SMO
- This takes two Lagrange multipliers α_i and α_j and adjusts them to maximise the dual objective function
- This is very quick as it can be done in closed form
- Note that because $\sum_{k=1}^m y_k \alpha_k = 0$ we have to change at least two variables at the same time
- A heuristic is used to choose good pairs of α 's to optimise
- Run until close to the optimum

Extended Feature Space

- We can generalise the SVM if we map all our features vectors to an extended feature space

$$x \rightarrow \phi(x)$$

- The components of $\phi(x)$ will typically be (non-linear) functions of x (e.g. $\phi_1(x) = x_1^2, \phi_2(x) = x_2^2, \phi_3(x) = \sqrt{2}x_1x_2$)
- We are free to choose whatever mappings we like
- There may be many more components of $\phi(x)$ than of x making it easier to find a linear separation of the two classes
- But in the extended feature space (involving $\phi(x) = (\phi_1(x), \phi_2(x), \dots, \phi_r(x))$) the time complexity is $O(r^3)$

Obtaining the Dual Form of the Problem

- Differentiating the Lagrangian with respect to w

$$\mathcal{L}(w, b, \alpha) = \frac{1}{2} \|w\|^2 - \sum_{k=1}^m \alpha_k (y_k (\langle w, \phi(x_k) \rangle - b) - 1)$$

- $\nabla_w \mathcal{L} = w - \sum_{k=1}^m \alpha_k y_k \phi(x_k) = 0$ implies that $w^* = \sum_{k=1}^m \alpha_k y_k \phi(x_k)$

- $\frac{\partial \mathcal{L}}{\partial b} = \sum_{k=1}^m \alpha_k y_k = 0$ implies $\sum_{k=1}^m \alpha_k y_k = 0$

- Substituting back into the Lagrangian

$$\max_{\alpha \geq 0} \sum_{k=1}^m \alpha_k - \frac{1}{2} \sum_{k,l=1}^m \alpha_k \alpha_l y_k y_l \langle \phi(x_k), \phi(x_l) \rangle$$

Kernel Trick

- We will show in the next lecture that if $K(x, y)$ is a positive semi-definite function then it can always be written as

$$K(x, y) = \langle \phi(x), \phi(y) \rangle$$


- As $\langle \phi(x_k), \phi(x_l) \rangle$ appears in the dual problem we can express the dual problem as finding α_k 's that maximise

$$\mathcal{L}(\alpha) = \sum_{k=1}^m \alpha_k - \frac{1}{2} \sum_{k,l=1}^m \alpha_k \alpha_l y_k y_l K(x_k, x_l)$$

- We therefore never have to compute $\phi(x)$

Soft Margins

- We can relax the margin constraints by introducing *slack variables*, $s_k \geq 0$



$$y_k(\langle x_k, w \rangle - b) \geq 1 - s_k$$

- Minimise $\frac{\|w\|^2}{2} + C \sum_{k=1}^n s_k$
- Larger C punishes slack variables more

Dual Problem with Slack Variables

- The Lagrangian with slack variables is

$$\mathcal{L} = \frac{1}{2}\|\mathbf{w}\|^2 + C \sum_{k=1}^m s_k - \sum_{k=1}^m \alpha_k \left(y_k (\langle \mathbf{w}, \phi(\mathbf{x}_k) \rangle - b) - 1 + s_k \right) - \sum_{k=1}^m \beta_k s_k$$

where β_k are Lagrange multipliers that ensure $s_k \geq 0$ (note that $\beta_k \geq 0$ —this is the KKT condition)


- Now minimising with respect to s_i

$$\frac{\partial \mathcal{L}}{\partial s_i} = C - \alpha_i - \beta_i = 0$$

- Or $\alpha_i = C - \beta_i$. Since $\beta_i \geq 0$ the constraint is $\alpha_i \leq C$ (recall also $\alpha_i \geq 0$)

Getting SVMs to Work Well

- SVMs rely on distances between data points
- These will change relative to each other if we rescale some features but not other—giving different maximum-margin hyper-planes



- If we don't know what features are important (most often the case), then it is worth scaling each feature (for example, so their range is between 0 and 1 or their variance is 1)

Choosing the Right Kernel Function

- There are kernels design for particular data types (e.g. string kernels for text or biological sequences)
- For numerical data, people tend to look at using no kernel (linear SVM), a radial basis function (Gaussian) kernel or polynomial kernels
- Kernels often come with parameters, e.g. the popular radial basis function kernel

$$K(\mathbf{x}, \mathbf{y}) = e^{-\gamma \|\mathbf{x} - \mathbf{y}\|^2}$$


- Optimal γ values range over 2^{-15} – 2^3

SVM Libraries

- Although SVMs have unique solutions, they require very well written optimisers
- If you have a large data set they can be very slow
- There are good libraries out there: svmlib, svm-lite, (now old), scikit-learn, etc.
- These will often automate normalisation of data and grid search for parameters

Outline

- The Big Picture
- Maximum Margins
- Duality
- Practice



Optimising C

- Recall that we can introduce soft-margins using slack variables where we minimise $\frac{\|\hat{\mathbf{w}}\|^2}{2} + C \sum_{k=1}^m s_k$ subject to constraints
- In practice it can make a huge difference to the performance if we change C
- Optimal C values changes by many orders of magnitude e.g. 2^{-5} – 2^{15}
- Typically optimised by a grid search (start from 2^{-5} say and double until you reach 2^{15})
- Measure performance on a validation set

Multi-Class Problems

- By construction SVMs separate only two classes
- If we have a multi-class problem we have to use multiple SVMs

- There are two major ways practitioners do this

One-versus-all: for each class, train a separate classifier to determine that class versus all others

All-pairs: train a classifier for all pairs of classes

- In both cases choose the class which the classifier is most certain about
- Beware SVMs don't like imbalanced datasets

Conclusions

- We've seen how SVMs work
- We've learnt how to use them
- We've seen that we can find the maximum margin hyper-plane by solving a quadratic programming problem (with a unique solution)
- This is a convex optimisation problem with a unique optimum
- The **dual problem** of an SVM is particularly simple, especially if we use a positive semi-definite kernel (we explore these in the next lecture)