
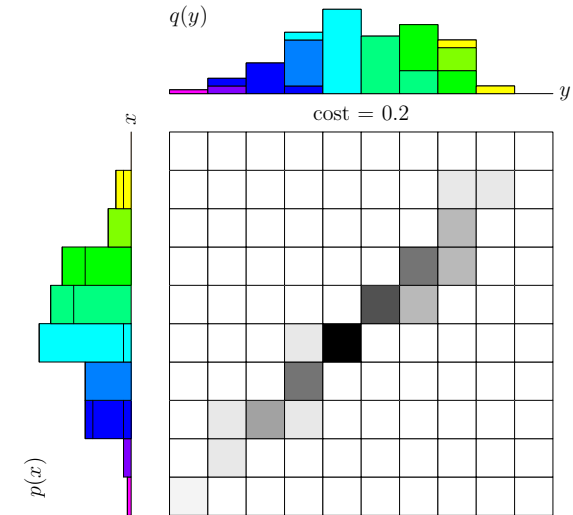


## Wasserstein GANs



GANs, Wasserstein distance, Duality, WGANs

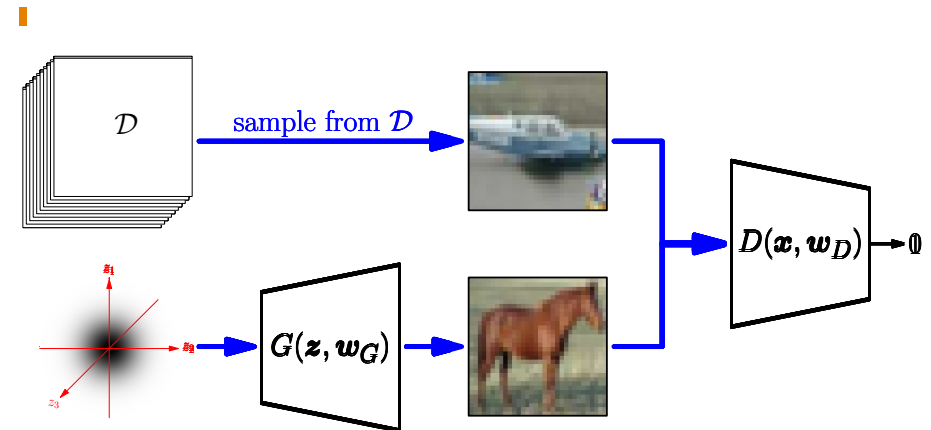
1. **GANs**
2. Wasserstein Distance
3. Wasserstein GANs



## Generative Adversarial Networks

- One of the applications of Deep Learning that has most excited the public are **Generative Adversarial Networks** or GANs
- Their aim is to generate new random samples from the same distribution as some training set,  $\mathcal{D}$
- Their number of real world applications are questionable
- But nobody cares because they are cool
- *Out of date warning:* someone invented diffusion models

## How GANs Work



# Training GANs

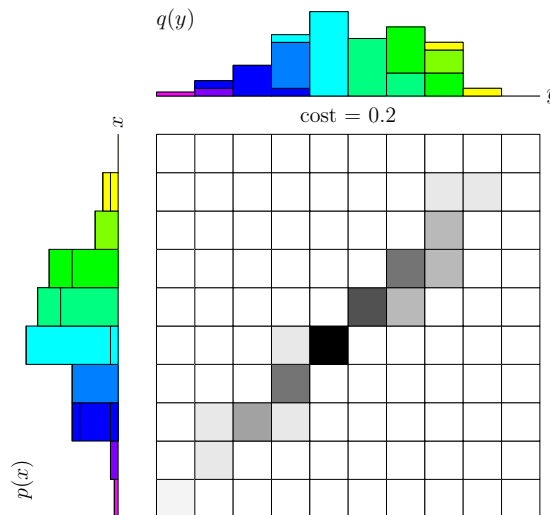
- The loss of the generator depends on its ability to trick the discriminator
- The loss of the discriminator depends on its ability not to be tricked
- We try to train the two networks simultaneously
- We hope that over time the generator produces better and better fakes

# Problems of GANs

- GANs are notoriously difficult to train
- The generator and discriminator training can decouple
- Often the discriminator becomes too good at correctly identifying the generated images
- Then there can be little gradient information to help the generator as every small change in parameters doesn't significantly change the discriminator decision
- To try to solve this problem we first make a seemingly unconnected diversion

## Outline

1. GANs
2. Wasserstein Distance
3. Wasserstein GANs



## Measuring Distances Between Distributions

- In many machine learning tasks we want to minimise the distance between two probability distributions
- This requires that we can measure distances between probability distributions
- One prominent measure is the Kullback-Leibler or KL divergence

$$\text{KL}(p||q) = \int p(x) \log \left( \frac{p(x)}{q(y)} \right) dx$$


- This is very commonly used in ML (e.g. VAEs, Variational Approximation)

## Trouble with KL

- KL-divergences are non-negative quantities that are minimised when the two probability distributions are the same
- They are not distances (they aren't symmetric and they don't satisfy the triangular inequality)

We don't really care about this, but what


- we do care about is that if  $q(x) = 0$  when  $p(x) \neq 0$  then  $\log\left(\frac{p(x)}{q(x)}\right)$  diverges
- We can therefore have distributies that seem very similar but their KL-divergence is huge (or infinite)



## Wasserstein Distance

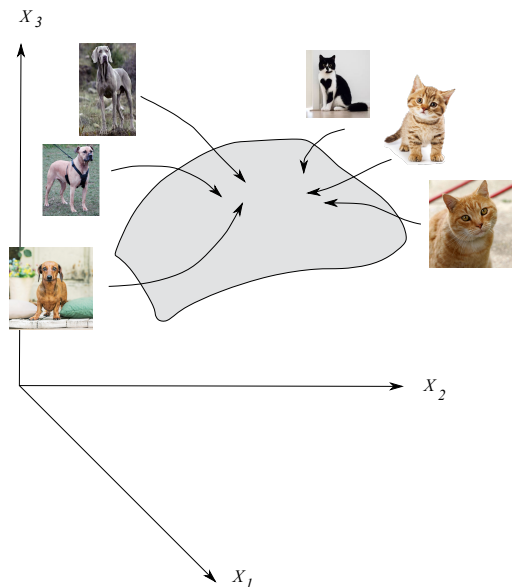
- A more benign measure of the differences between two probability functions is the **Wasserstein** or **Earth Moving** distance

This is a true distance, but more importantly for us it measure distance in a very natural way so that distributions that are close has a small Wasserstein distance



- Although this seems contrived if our probability distribution represents the probability of a  $128 \times 128$  matrix of real valued triples represents an image of dog, then it is easy to imagine that the Wasserstein distance may be more benign than the KL-divergence

## High Probability Manifold




## Transportation Policy

- But how do we formalise the Wasserstein distance?
- A good place to start is to define a transportation policy  $\gamma(x, y)$  with

$$\int \gamma(x, y) dy = p(x) \quad \int \gamma(x, y) dx = q(y)$$

- This looks like a joint probability distribution, but we interpret  $\gamma(x, y)$  as the amount of probability mass/density that we transfer from  $p(x)$  to  $q(y)$

## Transportation Policy



## The Cost of Transport


- We want to choose the transportation policy that minimises the amount of probability mass we need to move
- Let  $d(x, y) = \|x - y\|$  be a distance measure then the cost of a transportation policy is

$$C(\gamma) = \int \int d(x, y) \gamma(x, y) dx dy = \mathbb{E}_{\gamma}[d(x, y)]$$

where we interpret  $\gamma(x, y)$  as a probability distribution

- Usually we take  $d(x, y)$  to be the Euclidean distance, but we can choose any distance

## Transportation Cost



## The Wasserstein Distance

- The Wasserstein distance  $W(p, q)$  between two probability distributions is defined as

$$W(p, q) = \min_{\gamma \in \Lambda(p, q)} \mathbb{E}_{\gamma}[d(x, y)]$$

- Where  $\Lambda(p, q)$  is the set of joint distributions  $\gamma(x, y)$  such that

$$\int \gamma(x, y) dy = p(x) \quad \int \gamma(x, y) dx = q(y)$$

## Computing the Wasserstein Distance

- To compute the Wasserstein distance we have to solve a minimisation task!
- This looks nasty, but it is a (continuous) linear programming problem
- Suppose  $p$  and  $q$  were discrete distribution (i.e.  $x$  and  $y$  only take discrete points)
- Then we could treat each value of  $\gamma(x, y)$  as an element of a vector  $\gamma$  and each value of  $d(x, y)$  as an element of a vector  $D$
- Our objective is to choose  $\gamma$  to minimise  $D^T \gamma$

## Constraints

$$\sum_j \gamma(x_i, y_j) = p(x_i)$$

$$\sum_i \gamma(x_i, y_j) = q(y_j)$$

$$A\gamma = P$$

$$\begin{pmatrix} 1 & 1 & \cdots & 1 & 0 & 0 & \cdots & 0 & \cdots & 0 & 0 & 0 & \cdots & 0 \\ 0 & 0 & \cdots & 0 & 1 & 1 & \cdots & 1 & \cdots & 0 & 0 & 0 & \cdots & 0 \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots \\ 0 & 0 & \cdots & 0 & 0 & 0 & \cdots & 0 & \cdots & 1 & 1 & \cdots & 1 \\ 1 & 0 & \cdots & 0 & 1 & 0 & \cdots & 0 & \cdots & 1 & 0 & \cdots & 0 \\ 0 & 1 & \cdots & 0 & 0 & 1 & \cdots & 0 & \cdots & 0 & 1 & \cdots & 0 \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots \\ 0 & 0 & \cdots & 1 & 0 & 0 & \cdots & 1 & \cdots & 0 & 0 & \cdots & 1 \end{pmatrix} \begin{pmatrix} \gamma(x_1, y_1) \\ \gamma(x_2, y_1) \\ \vdots \\ \gamma(x_n, y_1) \\ \gamma(x_1, y_2) \\ \gamma(x_2, y_2) \\ \vdots \\ \gamma(x_n, y_2) \\ \vdots \\ \gamma(x_1, y_n) \\ \gamma(x_2, y_n) \\ \vdots \\ \gamma(x_n, y_n) \end{pmatrix} = \begin{pmatrix} q(y_1) \\ q(y_2) \\ \vdots \\ q(y_n) \\ p(x_1) \\ p(x_2) \\ \vdots \\ p(x_n) \end{pmatrix}$$

## Lagrange Formulation

- For discrete distributions

$$\min_{\gamma} D^T \gamma$$

subject to  $A\gamma = P, \quad \gamma \geq 0$

- Writing the Lagrangian

$$\mathcal{L}(\gamma, \alpha) = D^T \gamma - \alpha^T (A^T \gamma - P)$$

where  $\alpha$  is a vector of Lagrange multipliers

- The solution to the discrete optimisation problem is given by

$$\min_{\gamma} \max_{\alpha} \mathcal{L}(\gamma, \alpha)$$

## Dual Form

- We can rearrange

$$\begin{aligned} \mathcal{L}(\gamma, \alpha) &= D^T \gamma - \alpha^T (A^T \gamma - P) \\ &= P^T \alpha - \gamma^T (A^T \alpha - D) \end{aligned}$$

- We note that  $\gamma \geq 0$  so the dual problem is to find a vector  $\alpha$  that maximises  $P^T \alpha$  subject to the constraints  $A^T \alpha \leq D$
- Although the vector form allows us to make connections with our earlier discussion of linear programming, it is a little difficult to interpret

## Explicit Form

- We can write a Lagrangian for the original problem

$$\mathcal{L} = \sum_{i,j} d(\mathbf{x}_i, \mathbf{y}_j) \gamma(\mathbf{x}_i, \mathbf{y}_j) - \sum_i \alpha(\mathbf{x}_i) \left( \sum_j \gamma(\mathbf{x}_i, \mathbf{y}_j) - p(\mathbf{x}_i) \right) - \sum_j \beta(\mathbf{y}_j) \left( \sum_i \gamma(\mathbf{x}_i, \mathbf{y}_j) - q(\mathbf{y}_j) \right)$$

subject to  $\gamma(\mathbf{x}_i, \mathbf{y}_j) \geq 0$  where  $\alpha(\mathbf{x}_i)$  and  $\beta(\mathbf{y}_j)$  are Lagrange multipliers (they are components of  $\boldsymbol{\alpha}$ )

- Rearranging

$$\mathcal{L} = \sum_i \alpha(\mathbf{x}_i) p(\mathbf{x}_i) + \sum_j \beta(\mathbf{y}_j) q(\mathbf{y}_j) - \sum_{i,j} \gamma(\mathbf{x}_i, \mathbf{y}_j) (\alpha(\mathbf{x}_i) + \beta(\mathbf{y}_j) - d(\mathbf{x}_i, \mathbf{y}_j))$$

- This is equivalent to maximising  $\sum_i \alpha(\mathbf{x}_i) p(\mathbf{x}_i) + \sum_j \beta(\mathbf{y}_j) q(\mathbf{y}_j)$ , subject to

$$\forall i, j \quad \alpha(\mathbf{x}_i) + \beta(\mathbf{y}_j) \leq d(\mathbf{x}_i, \mathbf{y}_j)$$

## Dual Form Constraint

- We note that  $\alpha(\mathbf{x}) + \beta(\mathbf{y}) \leq d(\mathbf{x}, \mathbf{y})$  for all  $\mathbf{x}$  and  $\mathbf{y}$
- This has to be true when  $\mathbf{x} = \mathbf{y}$  so that

$$\alpha(\mathbf{x}) + \beta(\mathbf{x}) \leq d(\mathbf{x}, \mathbf{x}) = 0$$

- So  $\beta(\mathbf{x}) = -\alpha(\mathbf{x}) - \epsilon(\mathbf{x})$  where  $\epsilon(\mathbf{x}) \geq 0$
- But want to maximise

$$\int \alpha(\mathbf{x}) p(\mathbf{x}) d\mathbf{x} + \int \beta(\mathbf{y}) q(\mathbf{y}) d\mathbf{y} = \int \alpha(\mathbf{x}) (p(\mathbf{x}) - q(\mathbf{x})) d\mathbf{x} - \int q(\mathbf{x}) \epsilon(\mathbf{x}) d\mathbf{x}$$

- This is maximised when  $\epsilon(\mathbf{x}) = 0$  i.e.  $\beta(\mathbf{x}) = -\alpha(\mathbf{x})$

## Continuous Form

- We can write a Lagrangian for the continuous problem

$$\mathcal{L} = \iint d(\mathbf{x}, \mathbf{y}) \gamma(\mathbf{x}, \mathbf{y}) d\mathbf{x} d\mathbf{y} - \int \alpha(\mathbf{x}) \left( \int \gamma(\mathbf{x}, \mathbf{y}) d\mathbf{y} - p(\mathbf{x}) \right) d\mathbf{x} - \int \beta(\mathbf{y}) \left( \int \gamma(\mathbf{x}, \mathbf{y}) d\mathbf{x} - q(\mathbf{y}) \right) d\mathbf{y}$$

subject to  $\gamma(\mathbf{x}, \mathbf{y}) \geq 0$  where  $\alpha(\mathbf{x})$  and  $\beta(\mathbf{y})$  are Lagrange multiplier functions

- Rearranging

$$\mathcal{L} = \int \alpha(\mathbf{x}) p(\mathbf{x}) d\mathbf{x} + \int \beta(\mathbf{y}) q(\mathbf{y}) d\mathbf{y} - \iint \gamma(\mathbf{x}, \mathbf{y}) (\alpha(\mathbf{x}) + \beta(\mathbf{y}) - d(\mathbf{x}, \mathbf{y})) d\mathbf{x} d\mathbf{y}$$

- This is equivalent to maximising  $\int \alpha(\mathbf{x}) p(\mathbf{x}) d\mathbf{x} + \int \beta(\mathbf{y}) q(\mathbf{y}) d\mathbf{y}$ , subject to

$$\alpha(\mathbf{x}) + \beta(\mathbf{y}) \leq d(\mathbf{x}, \mathbf{y})$$

## Dual Form

- Thus the dual problem is to find a function  $\alpha(\mathbf{x})$ —or a vector of functions  $(\alpha(\mathbf{x}_i)|i)$ —that maximises

$$\int \alpha(\mathbf{x}) (p(\mathbf{x}) - q(\mathbf{x})) d\mathbf{x}$$

- Subject to the constraint

$$\alpha(\mathbf{x}) - \alpha(\mathbf{y}) \leq d(\mathbf{x}, \mathbf{y}) = \|\mathbf{x} - \mathbf{y}\|$$


- This is a continuity constraint on the Lagrange multiplier function  $\alpha(\mathbf{x})$  known as Lipschitz-1

## Lipschitz-1 Functions

- We note for a Lipschitz-1 function and any unit vector  $u$

$$u^T \nabla \alpha(x) = \lim_{\epsilon \rightarrow 0} \frac{\alpha(x) - \alpha(x + \epsilon u)}{\epsilon} \leq \frac{\epsilon}{\epsilon} = 1$$

- That is, at every point the gradient in all directions must be less than 1 (since the gradient defines the direction of greatest increase it is both necessary and sufficient for  $\|\nabla \alpha(x)\| \leq 1$  everywhere)



## Calculating the Wasserstein Distance

- To recall the big picture we want to compute the Wasserstein distance

$$W(p, q) = \min_{\gamma \in \Lambda(p, q)} \mathbb{E}_{\gamma}[d(x, y)]$$


- For high dimensional objects  $\gamma(x, y)$  would be a huge object to approximate
- Instead we can compute the Wasserstein distance in the dual formulation

$$W(p, q) = \max_{\alpha(x)} \int \alpha(x) (p(x) - q(x)) dx = \max_{\alpha} \mathbb{E}_p[\alpha(X)] - \mathbb{E}_q[\alpha(X)]$$

subject to the constraint that  $\alpha(x)$  is a Lipschitz-1 function

## Outline

1. GANs
2. Wasserstein Distance
3. Wasserstein GANs



## Back to GANs

- What has this got to do with GANs?
- Suppose we want to minimise the distance between the distribution  $p(x)$  of real images (of which  $\mathcal{D}$  are samples) and the distribution  $q(x)$  of images drawn from a generator
- We can use a normal GAN generator,  $G(z, w_G)$ , that generates an image when given a random variable  $z \sim \mathcal{N}(0, I)$
- To do this we choose the weights,  $w_G$  of the generator to minimise

$$W(p, q) = \max_{\alpha(x)} (\mathbb{E}_{x \sim p}[\alpha(x)] - \mathbb{E}_{x \sim q}[\alpha(x)])$$

## Estimating Expectations

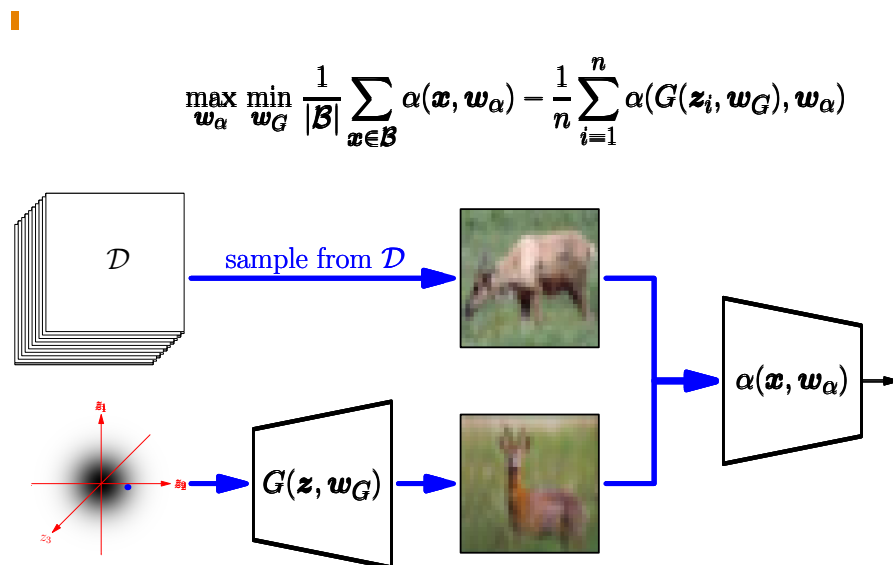
- Although we can't compute  $\mathbb{E}_p[\alpha(\mathbf{x})]$  and  $\mathbb{E}_q[\alpha(\mathbf{x})]$  exactly, we can estimate them from samples

$$\mathbb{E}_p[\alpha(\mathbf{x})] \approx \frac{1}{|\mathcal{B}|} \sum_{\mathbf{x} \in \mathcal{B}} \alpha(\mathbf{x}), \quad \mathbb{E}_q[\alpha(\mathbf{x})] \approx \frac{1}{n} \sum_{i=1}^n \alpha(G(\mathbf{z}_i, \mathbf{w}_G))$$

- where  $\mathcal{B} \subset \mathcal{D}$  is a minibatch of true images and  $\mathbf{z}_i \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
- From this we can choose  $\mathbf{w}_G$  to minimise

$$C = \frac{1}{|\mathcal{B}|} \sum_{\mathbf{x} \in \mathcal{B}} \alpha(\mathbf{x}) - \frac{1}{n} \sum_{i=1}^n \alpha(G(\mathbf{z}_i, \mathbf{w}_G))$$

## Wasserstein GANs



## The Critic

- For this quantity to approximate the Wasserstein distance we need to find a function  $\alpha(\mathbf{x}, \mathbf{w}_\alpha)$  that maximises  $C$
- To do this we learn a second network, the critic or discriminator whose job it is to maximise

$$C = \frac{1}{|\mathcal{B}|} \sum_{\mathbf{x} \in \mathcal{B}} \alpha(\mathbf{x}, \mathbf{w}_\alpha) - \frac{1}{n} \sum_{i=1}^n \alpha(G(\mathbf{z}_i, \mathbf{w}_G), \mathbf{w}_\alpha)$$

- The network  $\alpha(\mathbf{x}, \mathbf{w}_\alpha)$  should be Lipschitz-1 (which we usually botched by, for example, by setting the spectral norm of the convolutional weight matrix to 1)

## Lesson

- Wasserstein GANs are, at least for me, one of the most elegant pieces of theory in recent years
- By trying to minimise the Wasserstein distance between the distribution of a generator and a true distribution we arrive at optimising two adversarial networks just like a GAN
- This uses a rather beautiful dual formulation
- It is claimed that W-GANs solve many of the problems of traditional GANs