Advanced Machine Learning Subsidary Notes
Lecture 2: Over-Fitting

Adam Prügel-Bennett

January 28, 2021

1 Keywords

- Overfitting, regularisation, feature selection

2 Main Points

- Over-fitting

  - Over-fitting is when we learn **spurious rules** that explain the training set
  - If we use an infinitely flexible machine then we can't generalise at all

- Controlling Complexity

  - More training example improves generalisation by allowing us to eliminate spurious rules
  - To achieve good generalisation performance machine must match the underlying structure of the problem
  - By preprocessing or careful feature engineering you can sometimes make the learning task easier
  - Deep learning works because
    * It uses a lot of training data
    * They use convolutions that
      · have many fewer parameter to learn than fully connected layers
      · respect translation invariance
      · find local features (at different scales through the network)
  - We usually don't know the structure of the inputs (they are too high dimensional)
    * Try different machines to see what fits
    * Fine-tune the models (adjusting hyper-parameters)
    * Feature Engineering very often help
      · Feature selection
      · Using PCA or clustering
      · Normalise your input features
    * Need a validation set to do this
    * Beware you are very likely to over-estimate your generalisation error if you use your test set to select your model
    * Can use K-fold cross-validation to get a better estimate of generalisation error

- Regularisation

  - Adding regularisation terms can help choose a simpler model

- L_2 regularisers punish large weights making the fitting function smoother
- L_1 regularisers can do automatic feature selection
- Can do subtle forms of regularisation such as choosing a maximum margin solution

3   Exercises

3.1   Estimating the error in the mean

- Consider a binary classification problem
- Suppose we measure the accuracy on a validation set of size $n = 10\,000$
- If we correctly predict $8\,000$ of the validation set what is

  1. an estimate of the mean accuracy?
  2. the accuracy of this estimate?

- State you assumptions
- What would the answer be if $n = 100$ and we got 80 correct on the validation set?
- Answer at the end

4   Experiments

4.1   Overfitting Game

- Write your own simulation (see lecture notes)
- Generate n random numbers $X_i \sim U(0,3)$ (you can try different distributions)
- Add some noise to generate $Y_i = X_i + Z_i$ where $Z_i \sim \mathcal{N}(0,1)$
- Choose i^* with the largest value of Y_i
- Compute $\Delta = Z_{i^*} = Y_{i^*} - X_{i^*}$
- Repeat many times and plot a histogram

4.2   Cross-validation

- Try using cross validation

```python
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn import datasets
from sklearn import svm

X, y = datasets.load_iris(return_X_y=True)
X.shape, y.shape

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.4, random_state=0)

X_train.shape, y_train.shape

X_test.shape, y_test.shape
```

```
clf = svm.SVC(kernel='linear', C=1).fit(X_train, y_train)
clf.score(X_test, y_test)
```

4.3 Project

- There are a lot of practical tips that you can use in your project

    - Try them out and see what works and what doesn't
    - Many modifications won't make much difference, but occasionally one will make a lot of difference
    - Be very wary of estimating your generalisation performance
        * You are measuring on a finite validation set so you will have errors
        * Once you've used your validation set for selecting models, your likely to over-estimate your generalisation error

5 Answers

5.1 Estimating the error in the mean

- We assume that the examples in the validation set are independent so the number of successes, k, will be binomially distributed

$$\text{Bin}(k|n, p) = \binom{n}{k} p^k (1-p)^{n-k}$$

- We don't know the success probability p, but an unbiased estimate of it is the number of observed success divided by n, that is 0.8

- For a binomial the variance in k is given by $n\,p\,(1-p)$ or 1600 if $p = 0.8$ and $n = 10\,000$

- The typical size of the fluctuations in k would be $\sqrt{n\,p\,(1-p)} = 40$

- Our estimate of p would therefore typically have fluctuations of size $40/n = 0.004$

- Thus our estimated success rate is 0.8 ± 0.004 or 80% with an error of about half a percent

- This is small because we used a large validation set. If $n = 100$ and we got 80 successes then we would have an error of 0.8 ± 0.04 so we have an error of 4%.

- Note that by choosing the best of 10 machines we might expect to over-estimate the performance by around 5% in this case